



## Contents

---

# Introduction to Linear Regression

## solidstar Why This Algorithm Matters

Simple Linear Regression is a **super important algorithm** because the techniques that we are going to learn in this will also be applicable in **deep learning**. When you're learning the first neural network, which is called an **Artificial Neural Network (ANN)**, the same foundational concepts apply.

## Problem Statement

### solidquestion-circle What Problems Does Linear Regression Solve?

As the name suggests — **Regression**. In supervised machine learning technique, if you have a **regression problem statement**, we can solve it with the help of simple linear regression.

## Understanding Through Example

Let's consider a specific and easy dataset:


Weight (kg)	Height (cm)
74	170
80	180
75	175.5

Table 1: Sample Dataset for Linear Regression

## solidbullseye Main Aim

Our main aim is that we need to **train a model**. This specific model, whenever we give a **new weight**, should be able to **predict the height**.

## Feature Types




## Simple vs Multiple Linear Regression

### Key Distinction

- **Simple Linear Regression:** When we have just **ONE input feature** and one output feature
- **Multiple Linear Regression:** When we have **multiple input features**

It is very important to understand simple linear regression first. Then automatically whatever terminologies and maths you learn here will get applied to multiple linear regression also.

## Geometric Interpretation: The Best Fit Line



### The Core Idea

With the help of simple linear regression, we try to create a **best fit line**. This best fit line will help us to do the prediction for the new weight.

#### How does prediction happen?

1. Once we get the best fit line
2. When we get a new data point (new weight)
3. We project it onto the best fit line
4. Read the corresponding  $y$ -value (height)

The best fit line should be created such that the **summation of all distances** (errors) between true points and predicted points should be **minimal**.

## Mathematical Notation & Fundamentals

### Important Note

Before we understand the maths behind it, we really need to understand some of the notation that will be used while explaining this entire machine learning algorithm.

## Equation of the Best Fit Line


### Different Notations for the Same Equation

The equation of a straight line can be written in multiple ways:

Notation	Form
Standard Form	$y = mx + c$
Research Paper Style	$y = \beta_0 + \beta_1 x$
Andrew Ng's Notation	$h_{\theta}(x) = \theta_0 + \theta_1 x$

We will use Andrew Ng's notation:  $h_{\theta}(x) = \theta_0 + \theta_1 x$

## Understanding Each Component



## Understanding $\theta_0$ (Intercept)

### **solidmap-marker-alt** The Intercept $\theta_0$


$\theta_0$  is the Intercept

If my  $x$  value is zero, then:

$$h_{\theta}(x) = \theta_0 + \theta_1 \cdot 0 = \theta_0$$

This indicates: When  $x$  is zero, where does the line meet the  $y$ -axis?

The point where the line intersects the  $y$ -axis is called the **intercept**.




## Understanding $\theta_1$ (Slope/Coefficient)

### **solidchart-line** The Slope $\theta_1$

$\theta_1$  is the Slope or Coefficient

It indicates: With unit movement in the  $x$ -axis, what is the movement with respect to the  $y$ -axis?

This is called the **slope movement**.



## Extension to Multiple Features

### Multiple Independent Features

If we have many independent features, the equation becomes:

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3 + \dots + \theta_n x_n$$

- $\theta_0$  — Always ONE intercept
- $\theta_1, \theta_2, \dots, \theta_n$  — Slopes for each feature
- $x_1, x_2, \dots, x_n$  — Independent features

Since we have ONE independent feature in simple linear regression, we have just ONE slope ( $\theta_1$ ).

## Predicted Points Notation


### The $\hat{y}$ Notation

The predicted points are denoted by:

$$\hat{y} = h_{\theta}(x)$$

Where  $\hat{y}$  (“y-hat”) represents the **predicted points** obtained from our model.

## Understanding Error



### Error Definition

$$\text{Error} = y - \hat{y}$$

**Error** is the difference between:

- $y$  — The true point (actual output from data)

- $\hat{y}$  — The predicted point (from best fit line)

### solidbullseye Main Aim Restated

Our main aim is to create a **best fit line** wherein when we try to calculate or do the **summation of all errors**, it should be **minimal**.

If it is not minimal and there is another best fit line that minimizes this error, we select that specific best fit line instead.

# Cost Function: Mean Squared Error

## The Need for a Cost Function

### solidexclamation-circle The Problem

We cannot just randomly create many best fit lines and check which one has the least error. This is **not efficient**.

We need to find an **optimized way**:

1. Create one best fit line
2. Rotate this line by changing values of intercept ( $\theta_0$ ) and coefficient ( $\theta_1$ )
3. Find the best fit line with **minimal error**

## The Cost Function Formula

### solidfunction Mean Squared Error (MSE) Cost Function

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2 \quad (1)$$

Where:

- $J(\theta_0, \theta_1)$  — Cost Function
- $m$  — Number of data points
- $h_\theta(x^{(i)})$  — Predicted points
- $y^{(i)}$  — True points (actual output)
- The squaring makes this the “Mean **Squared** Error”

### solidquestion-circle Why Squared?

The reason we are **squaring** is because the technique/cost function we are using is called **Mean Squared Error (MSE)**. There are advantages of using MSE which we will discuss. Other cost functions exist:

- Mean Absolute Error (MAE)
- Root Mean Squared Error (RMSE)

## Our Final Objective

solidflag-checked What We Need to Solve

**MINIMIZE**  $J(\theta_0, \theta_1)$

$$\min_{\theta_0, \theta_1} J(\theta_0, \theta_1) = \min_{\theta_0, \theta_1} \frac{1}{2m} \sum_{i=1}^m \left( h_\theta(x^{(i)}) - y^{(i)} \right)^2$$

**How?** By continuously changing  $\theta_0$  and  $\theta_1$  values to find the best fit line where the summation of error is **minimal as possible**.

## Gradient Descent: Finding the Best Fit Line

### Simplified Example

To understand gradient descent in 2D, let's assume  $\theta_0 = 0$  (line passes through origin).


#### soliddatabase Sample Dataset

Consider the dataset:

x	y
1	1
2	2
3	3

With  $\theta_0 = 0$ , our equation becomes:  $h_\theta(x) = \theta_1 \cdot x$

### Case 1: $\theta_1 = 1$



#### solidsquare-root-alt Important Formula


**Computing Cost Function when  $\theta_1 = 1$ :**

$$J(\theta_1) = \frac{1}{2 \cdot 3} [(1 - 1)^2 + (2 - 2)^2 + (3 - 3)^2] \quad (2)$$

$$= \frac{1}{6} [0 + 0 + 0] \quad (3)$$

$$= [0] \quad (4)$$

**Result:** Cost function is **ZERO!** The line passes through ALL points perfectly.

**Case 2:  $\theta_1 = 0.5$** **solidsquare-root-alt Important Formula****Computing Cost Function when  $\theta_1 = 0.5$ :**

$$J(\theta_1) = \frac{1}{2 \cdot 3} [(0.5 - 1)^2 + (1 - 2)^2 + (1.5 - 3)^2] \quad (5)$$

$$= \frac{1}{6} [0.25 + 1 + 2.25] \quad (6)$$

$$= \frac{3.5}{6} \approx [0.58] \quad (7)$$

**Case 3:  $\theta_1 = 0$** **solidsquare-root-alt Important Formula****Computing Cost Function when  $\theta_1 = 0$ :**


$$J(\theta_1) = \frac{1}{2 \cdot 3} [(0 - 1)^2 + (0 - 2)^2 + (0 - 3)^2] \quad (8)$$

$$= \frac{1}{6} [1 + 4 + 9] \quad (9)$$

$$= \frac{14}{6} \approx [2.33] \quad (10)$$

This is NOT a good fit! Error is quite high.

## The Gradient Descent Curve




### solidmountain The Gradient Descent Curve

This entire curve is called the **Gradient Descent** curve.

- The lowest point is called the **Global Minima**
- At the global minima, our cost function is minimal
- This is where we get our **best fit line**
- We need to reach this point by changing  $\theta_0$  and  $\theta_1$  values

**Key Insight:** Gradient descent is super important even in deep learning techniques!

## 3D Visualization (When $\theta_0 \neq 0$ )



**solidexclamation-triangle Important Note**

When we have both  $\theta_0$  and  $\theta_1$ , the gradient descent becomes a **3D surface** (like an inverted mountain or bowl). Both parameters must converge to their optimal values to reach the global minima.

# The Convergence Algorithm

## The Problem with Random Selection

### solidtimes-circle Inefficient Approach

When we were finding our cost function, we were **randomly changing** our  $\theta_1$  value (slope):

- First time:  $\theta_1 = 1$
- Second time:  $\theta_1 = 0.5$
- Third time:  $\theta_1 = 0$

This is **NOT an efficient technique!** We need an algorithm that:

1. Initializes one  $\theta_1$  value
2. Automatically increases or decreases based on the gradient descent

# The Convergence Algorithm

### solidrepeat The Convergence Algorithm

#### Repeat until convergence:

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1) \quad (11)$$

Where:

- $\theta_j$  — The parameter being updated
- $\alpha$  — Learning rate (small positive number)
- $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$  — Derivative (slope) of cost function
- $:=$  — Assignment operator (iterative update)

“Until convergence” means until we reach the global minima point.

## Understanding the Derivative (Slope)

### solidcalculator What Does the Derivative Tell Us?

The derivative  $\frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$  calculates the **slope** at the current point on the gradient descent curve.


We create a **tangent line** at the current point and determine:

- Is it a **positive slope** or **negative slope**?

**How to identify:** Look at the **right side** of the tangent line:

- If pointing **upward** → Positive slope
- If pointing **downward** → Negative slope

### Case 1: Negative Slope (Left of Minima)




### solidarrow-right When Slope is Negative

$$\theta_j := \theta_j - \alpha \cdot (\text{negative value})$$

$$\theta_j := \theta_j + (\text{positive value})$$

**Result:**  $\theta_j$  **INCREASES** → Moves toward global minima!

## Case 2: Positive Slope (Right of Minima)



**solidarrow-left When Slope is Positive**

$$\theta_j := \theta_j - \alpha \cdot (\text{positive value})$$

$$\theta_j := \theta_j - (\text{positive value})$$

**Result:**  $\theta_j$  DECREASES → Moves toward global minima!

## The Learning Rate ( $\alpha$ )

**solidtachometer-alt Learning Rate  $\alpha$**

**Learning Rate ( $\alpha$ )** is a smaller value that controls the **speed of convergence**.

- Typically:  $\alpha = 0.001$  (common practice)
- In sklearn library:  $\alpha = 0.001$  is the default for simple linear regression

**$\alpha$  too small**

Takes more time  
to converge



**$\alpha$  just right**

Optimal  
convergence



**$\alpha$  too big**

May never converge  
(jumps around)



**soliduser-tie Interview Question**

**Q: What is the importance of learning rate?**

**A:** The learning rate controls the convergence rate — how slow or fast the convergence should happen. A very small value takes more time; a very big value may cause the algorithm to never converge (continuously jumping).

## Complete Convergence Equations

### The Full Update Equations

#### solidcode Final Convergence Algorithm

**Repeat until convergence:**

**Update  $\theta_0$  (Intercept):**

$$\theta_0 := \theta_0 - \alpha \cdot \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})$$

**Update  $\theta_1$  (Slope):**

$$\theta_1 := \theta_1 - \alpha \cdot \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

**Where:**

- $\alpha$  = Learning rate
- $m$  = Number of data points
- $h_\theta(x^{(i)}) = \theta_0 + \theta_1 x^{(i)}$
- $y^{(i)}$  = Actual output

## Derivation Summary

#### solidgraduation-cap How We Got These Equations

Starting from:

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)})^2$$

**For  $j = 0$  (intercept):**


$$\frac{\partial}{\partial \theta_0} J = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot 1$$

The derivative of  $(\theta_0 + \theta_1 x)$  w.r.t.  $\theta_0$  is 1.

**For  $j = 1$  (slope):**


$$\frac{\partial}{\partial \theta_1} J = \frac{1}{m} \sum_{i=1}^m (h_\theta(x^{(i)}) - y^{(i)}) \cdot x^{(i)}$$

The derivative of  $(\theta_0 + \theta_1 x)$  w.r.t.  $\theta_1$  is  $x$ .



## Multiple Linear Regression

### Simple vs Multiple Linear Regression



### House Pricing Example

solidhome House Pricing Dataset

No. of Rooms	Size of House	Location	Price
$x_1$	$x_2$	$x_3$	$y$
<b>Independent Features</b>	<b>Dependent</b>		

Here we have **3 input features**, so this is **Multiple Linear Regression**.

### The Extended Equation

solidsuperscript Multiple Linear Regression Equation

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \theta_3 x_3$$

Term	Meaning	Example
$\theta_0$	Intercept (always ONE)	Y-intercept
$\theta_1$	Coefficient for $x_1$	Slope for “Number of Rooms”
$\theta_2$	Coefficient for $x_2$	Slope for “Size of House”
$\theta_3$	Coefficient for $x_3$	Slope for “Location”

**Key Point:** As the number of input features increases, that many coefficients will increase. But  $\theta_0$  (intercept) is always **ONE**.

## Gradient Descent in Higher Dimensions

### solidcubes Higher Dimensional Gradient Descent

With multiple features:

- The gradient descent surface becomes multi-dimensional
- Cannot draw 4D+ diagrams, but imagine an **inverted mountain**
- All coefficients ( $\theta_0, \theta_1, \theta_2, \theta_3, \dots$ ) must converge to the global minima

The cost function becomes:

$$J(\theta_0, \theta_1, \theta_2, \dots, \theta_n) = \frac{1}{2m} \sum_{i=1}^m \left( h_\theta(x^{(i)}) - y^{(i)} \right)^2$$

## Performance Metrics: $R^2$ and Adjusted $R^2$

### R-Squared ( $R^2$ )

**R-Squared Formula**


$$R^2 = 1 - \frac{SS_{\text{residual}}}{SS_{\text{total}}}$$

$$R^2 = 1 - \frac{\sum_{i=1}^m (y_i - \hat{y}_i)^2}{\sum_{i=1}^m (y_i - \bar{y})^2}$$

Where:

- $SS_{\text{residual}}$  = Sum of Squared Residuals (errors)
- $SS_{\text{total}}$  = Sum of Squared Total (from mean)
- $y_i$  = True values
- $\hat{y}_i$  = Predicted values
- $\bar{y}$  = Mean of true values

### Visual Understanding of $R^2$



**Interpreting  $R^2$**



- $R^2 = 0.70$  means 70% accuracy

- $R^2 = 0.85$  means 85% accuracy
- $R^2 = 0.90$  means 90% accuracy
- Higher  $R^2$  (closer to 1) → More accurate model

## The Problem with $R^2$

### solidbug Problem with R-Squared

Even if you add a feature that is **NOT correlated** with the output feature, the  $R^2$  value will still **increase**!

**Example:**

Features Used	$R^2$ Value
Size of House	0.75
Size + No. of Bedrooms	0.80
Size + Bedrooms + Location	0.85
Size + Bedrooms + Location + <b>Gender (irrelevant!)</b>	0.87

Gender has **NO correlation** with price, yet  $R^2$  still increased!

## Adjusted R-Squared

### solidbalance-scale Adjusted R-Squared Formula

$$R^2_{\text{adjusted}} = 1 - (1 - R^2) \cdot \frac{n - 1}{n - p - 1}$$

Where:

- $n$  = Number of data points
- $p$  = Number of independent features
- $R^2$  = Original R-squared value

### solidmagic How Adjusted $R^2$ Works

Adjusted  $R^2$  **penalizes** for adding features that are **not correlated** with the output.

Scenario	$R^2$	Adjusted $R^2$
2 correlated features	90%	86%
Add 1 correlated feature	92%	<b>88%</b> (increases)
Add 1 <b>uncorrelated</b> feature	93%	<b>83%</b> (decreases!)

### Key Insight:

- If new feature is correlated → Adjusted  $R^2$  **increases**

- If new feature is NOT correlated → Adjusted  $R^2$  **decreases**

# Complete Summary

## solidclipboard-check Summary

### Key Concepts Covered:

1. **Simple Linear Regression:** ONE input feature, creates best fit line

2. **Equation:**  $h_{\theta}(x) = \theta_0 + \theta_1 x$

- $\theta_0$  = Intercept (where line meets y-axis)
- $\theta_1$  = Slope (rate of change)

3. **Cost Function (MSE):**

$$J(\theta_0, \theta_1) = \frac{1}{2m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

4. **Gradient Descent:** Bowl-shaped curve; aim is to reach **Global Minima**

5. **Convergence Algorithm:**

$$\theta_j := \theta_j - \alpha \cdot \frac{\partial}{\partial \theta_j} J(\theta_0, \theta_1)$$

6. **Learning Rate ( $\alpha$ ):** Controls speed of convergence (typically 0.001)

7. **Multiple Linear Regression:** Multiple input features

$$h_{\theta}(x) = \theta_0 + \theta_1 x_1 + \theta_2 x_2 + \dots + \theta_n x_n$$

8. **Performance Metrics:**

- $R^2$  = Coefficient of determination (can be misleading)
- Adjusted  $R^2$  = Penalizes uncorrelated features (more reliable)

solidcheck-circle These techniques are foundational for Deep Learning & Neural Networks!

## End of Linear Regression Notes

solidbrain Foundation for Machine Learning & Deep Learning

solidrocket Ready for Practical Implementation!